

TensorPRO (Tensorflow Privacy Remindful Optimization)

Alessio Proietti IN550 Final Exam

Abstract:

L'idea di base è capire se un cliente in un determinato contesto socioeconomico contattato dalla banca sottoscriverà o no un deposito. Il task è di apprendimento supervisionato, le label sono nella colonna 'y'.

Il dataset <https://archive.ics.uci.edu/ml/datasets/Bank+Marketing> (<https://archive.ics.uci.edu/ml/datasets/Bank+Marketing>) è stato preliminarmente esplorato. In una seconda fase è stata standardizzata ogni feature numerica, quelle categoriali sono state codificate con la strategia one hot encoding.

Il dataset era fortemente sbilanciato, nuove istanze per l'allenamento sono state generate con l'algoritmo ADASYN. In conclusione è stata allenata una rete con ottimizzazione ADAM in modalità differential privacy e si sono confrontate delle metriche con la versione non differential private di ADAM.

Fase Esplorativa

In [1]:

```
# importo alcune libraries di cui avrò bisogno fin dall' inizio
import pandas as pd
import numpy as np

# libs per visualizzazione
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(color_codes=True)
```

In [2]:

```
df = pd.read_csv('https://raw.githubusercontent.com/alessio-proietti/2021-IN550-EXAM/main/data.csv', sep=';')
df.head(5)
```

Out[2]:

| | age | job | marital | education | default | housing | loan | contact | month | day_of_w |
|---|-----|-----------|---------|-------------|---------|---------|------|-----------|-------|----------|
| 0 | 56 | housemaid | married | basic.4y | no | no | no | telephone | may | 1 |
| 1 | 57 | services | married | high.school | unknown | no | no | telephone | may | 1 |
| 2 | 37 | services | married | high.school | no | yes | no | telephone | may | 1 |
| 3 | 40 | admin. | married | basic.6y | no | no | no | telephone | may | 1 |
| 4 | 56 | services | married | high.school | no | no | yes | telephone | may | 1 |

5 rows × 21 columns

In [3]:

```
df.tail(5)
```

Out[3]:

| | age | job | marital | education | default | housing | loan | contact | month | d |
|-------|-----|-------------|---------|---------------------|---------|---------|------|----------|-------|---|
| 41183 | 73 | retired | married | professional.course | no | yes | no | cellular | nov | |
| 41184 | 46 | blue-collar | married | professional.course | no | no | no | cellular | nov | |
| 41185 | 56 | retired | married | university.degree | no | yes | no | cellular | nov | |
| 41186 | 44 | technician | married | professional.course | no | no | no | cellular | nov | |
| 41187 | 74 | retired | married | professional.course | no | yes | no | cellular | nov | |

5 rows × 21 columns

In [4]:

```
# Con questo posso avere un quadro colonne non numeriche
df.describe(include = 'object')
```

Out[4]:

| | job | marital | education | default | housing | loan | contact | month | day_of_w |
|--------|--------|---------|-------------------|---------|---------|-------|----------|-------|----------|
| count | 41188 | 41188 | 41188 | 41188 | 41188 | 41188 | 41188 | 41188 | 41 |
| unique | 12 | 4 | 8 | 3 | 3 | 3 | 2 | 10 | |
| top | admin. | married | university.degree | no | yes | no | cellular | may | |
| freq | 10422 | 24928 | 12168 | 32588 | 21576 | 33950 | 26144 | 13769 | 8 |

In [5]:

```
# Posso estrarre informazioni di base sulle colonne numeriche
df.describe()
```

Out[5]:

| | age | duration | campaign | pdays | previous | emp.var.rate | c |
|-------|-------------|--------------|--------------|--------------|--------------|--------------|---|
| count | 41188.00000 | 41188.000000 | 41188.000000 | 41188.000000 | 41188.000000 | 41188.000000 | . |
| mean | 40.02406 | 258.285010 | 2.567593 | 962.475454 | 0.172963 | 0.081886 | |
| std | 10.42125 | 259.279249 | 2.770014 | 186.910907 | 0.494901 | 1.570960 | |
| min | 17.00000 | 0.000000 | 1.000000 | 0.000000 | 0.000000 | -3.400000 | |
| 25% | 32.00000 | 102.000000 | 1.000000 | 999.000000 | 0.000000 | -1.800000 | |
| 50% | 38.00000 | 180.000000 | 2.000000 | 999.000000 | 0.000000 | 1.100000 | |
| 75% | 47.00000 | 319.000000 | 3.000000 | 999.000000 | 0.000000 | 1.400000 | |
| max | 98.00000 | 4918.000000 | 56.000000 | 999.000000 | 7.000000 | 1.400000 | |

In [6]:

```
# Stampo i tipi di tutte le variabili per capire con cosa ho a che fare
df.dtypes
```

Out[6]:

```
age                int64
job                object
marital            object
education          object
default            object
housing            object
loan               object
contact            object
month              object
day_of_week        object
duration           int64
campaign           int64
pdays             int64
previous           int64
poutcome           object
emp.var.rate       float64
cons.price.idx     float64
cons.conf.idx      float64
euribor3m          float64
nr.employed        float64
y                  object
dtype: object
```

In [7]:

```
# Voglio contare i campi nulli o comunque capire se ce ne sono  
print(df.isnull().sum())
```


```
age          0  
job          0  
marital      0  
education    0  
default      0  
housing      0  
loan         0  
contact      0  
month        0  
day_of_week  0  
duration     0  
campaign     0  
pdays       0  
previous     0  
poutcome     0  
emp.var.rate 0  
cons.price.idx 0  
cons.conf.idx 0  
euribor3m    0  
nr.employed  0  
y            0  
dtype: int64
```

In [8]:

```
# È interessante calcolare e avere una visione delle correlazione tra le variabili numeriche
plt.figure(figsize=(20,10))
c = df.corr()
sns.heatmap(c,cmap="BrBG",annot=True)
c
```

Out[8]:


| | age | duration | campaign | pdays | previous | emp.var.rate | cons.price |
|----------------|-----------|-----------|-----------|-----------|-----------|--------------|------------|
| age | 1.000000 | -0.000866 | 0.004594 | -0.034369 | 0.024365 | -0.000371 | 0.000 |
| duration | -0.000866 | 1.000000 | -0.071699 | -0.047577 | 0.020640 | -0.027968 | 0.005 |
| campaign | 0.004594 | -0.071699 | 1.000000 | 0.052584 | -0.079141 | 0.150754 | 0.127 |
| pdays | -0.034369 | -0.047577 | 0.052584 | 1.000000 | -0.587514 | 0.271004 | 0.078 |
| previous | 0.024365 | 0.020640 | -0.079141 | -0.587514 | 1.000000 | -0.420489 | -0.203 |
| emp.var.rate | -0.000371 | -0.027968 | 0.150754 | 0.271004 | -0.420489 | 1.000000 | 0.775 |
| cons.price.idx | 0.000857 | 0.005312 | 0.127836 | 0.078889 | -0.203130 | 0.775334 | 1.000 |
| cons.conf.idx | 0.129372 | -0.008173 | -0.013733 | -0.091342 | -0.050936 | 0.196041 | 0.058 |
| euribor3m | 0.010767 | -0.032897 | 0.135133 | 0.296899 | -0.454494 | 0.972245 | 0.688 |
| nr.employed | -0.017725 | -0.044703 | 0.144095 | 0.372605 | -0.501333 | 0.906970 | 0.522 |



In [9]:


```
# Voglio vedere se la distribuzione dei contatti sulla settimana è uniforme  
# TL;DR lo è in buona approssimazione
```

```
df['day_of_week'].hist().plot(kind="bar", figsize=(10,5))  
plt.title("Numero di contatti v Giorno della settimana")  
plt.ylabel('Giorno della settimana')  
plt.xlabel('Numero di contatti');
```



In [10]:

```
# Come sono distribuiti contattati per età?  
df.age.hist().plot(kind="bar", figsize=(10,5))  
plt.title("Distribuzione contattati per età")  
plt.ylabel('Numero contatti')  
plt.xlabel('Età');
```




In [11]:

```
# Il dataset è sbilanciato, ecco un grafico a torta per capire le proporzioni  
df.y.value_counts().plot(kind='pie', subplots=True, startangle=90,  
figsize=(10,5), autopct='%1.1f%%')
```

Out[11]:

```
array([<matplotlib.axes._subplots.AxesSubplot object at 0x7f94c05fa0  
40>],  
      dtype=object)
```




In [12]:

```
# Voglio vedere come cambia il rapporto YES-NO in base al lavoro  
# Non sembra in effetti cambiare  
  
df.groupby('job').y.value_counts().plot(kind='bar')
```

Out[12]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f94c05c53d0>



In [13]:

```
# Stampa scatterplot per le variabili numeriche
sns.pairplot(df, corner=True)

# non sembrano emergere dei trend
```



Preparazione del Dataset

In [14]:

```
# separo le label dal resto dei dati
labels = df[['y']]
df.drop(columns=['y'])
```

Out[14]:

| | age | job | marital | education | default | housing | loan | contact | month |
|-------|-----|-------------|---------|---------------------|---------|---------|------|-----------|-------|
| 0 | 56 | housemaid | married | basic.4y | no | no | no | telephone | may |
| 1 | 57 | services | married | high.school | unknown | no | no | telephone | may |
| 2 | 37 | services | married | high.school | no | yes | no | telephone | may |
| 3 | 40 | admin. | married | basic.6y | no | no | no | telephone | may |
| 4 | 56 | services | married | high.school | no | no | yes | telephone | may |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 41183 | 73 | retired | married | professional.course | no | yes | no | cellular | no |
| 41184 | 46 | blue-collar | married | professional.course | no | no | no | cellular | no |
| 41185 | 56 | retired | married | university.degree | no | yes | no | cellular | no |
| 41186 | 44 | technician | married | professional.course | no | no | no | cellular | no |
| 41187 | 74 | retired | married | professional.course | no | yes | no | cellular | no |

41188 rows × 20 columns

In [15]:

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

In [16]:

```
# inizializzo gli oggetti che mi permetteranno di codificare le colonne categori
ali
le = LabelEncoder()
ohe = OneHotEncoder()
```

In [17]:

```
# estraggo le variabili categoriali e mostro un sample, la testa
categorical = df.select_dtypes(include=[object])
categorical.head(6)
```

Out[17]:

| | job | marital | education | default | housing | loan | contact | month | day_of_week |
|---|-----------|---------|-------------|---------|---------|------|-----------|-------|-------------|
| 0 | housemaid | married | basic.4y | no | no | no | telephone | may | mon |
| 1 | services | married | high.school | unknown | no | no | telephone | may | mon |
| 2 | services | married | high.school | no | yes | no | telephone | may | mon |
| 3 | admin. | married | basic.6y | no | no | no | telephone | may | mon |
| 4 | services | married | high.school | no | no | yes | telephone | may | mon |
| 5 | services | married | basic.9y | unknown | no | no | telephone | may | mon |

In [18]:

```
# One Hot Encoding delle variabili categoriali
categorical_le = categorical.apply(le.fit_transform)
categorical_sparse = ohe.fit_transform(categorical_le).toarray()

categorical_encoded = pd.DataFrame(categorical_sparse)
categorical_encoded
```

Out[18]:

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | ... | 45 | 46 | 47 | 48 | 49 | 50 | 51 | 52 |
|-------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 1 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | ... | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 2 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | ... | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 3 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 4 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | ... | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 41183 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 41184 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 41185 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 41186 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | ... | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 41187 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |

41188 rows × 55 columns

In [19]:

```
# Estraggo le variabili numeriche
numerical_not_scaled = df.select_dtypes(include=['float64', 'int64'])
numerical_not_scaled
```

Out[19]:

| | age | duration | campaign | pdays | previous | emp.var.rate | cons.price.idx | cons.conf.idx |
|-------|-----|----------|----------|-------|----------|--------------|----------------|---------------|
| 0 | 56 | 261 | 1 | 999 | 0 | 1.1 | 93.994 | -36.4 |
| 1 | 57 | 149 | 1 | 999 | 0 | 1.1 | 93.994 | -36.4 |
| 2 | 37 | 226 | 1 | 999 | 0 | 1.1 | 93.994 | -36.4 |
| 3 | 40 | 151 | 1 | 999 | 0 | 1.1 | 93.994 | -36.4 |
| 4 | 56 | 307 | 1 | 999 | 0 | 1.1 | 93.994 | -36.4 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 41183 | 73 | 334 | 1 | 999 | 0 | -1.1 | 94.767 | -50.8 |
| 41184 | 46 | 383 | 1 | 999 | 0 | -1.1 | 94.767 | -50.8 |
| 41185 | 56 | 189 | 2 | 999 | 0 | -1.1 | 94.767 | -50.8 |
| 41186 | 44 | 442 | 1 | 999 | 0 | -1.1 | 94.767 | -50.8 |
| 41187 | 74 | 239 | 3 | 999 | 1 | -1.1 | 94.767 | -50.8 |

41188 rows × 10 columns

In [20]:

```
# Per guadagnare performance apporterò una standardizzazione ai dati numerici
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```


In [21]:

```
numerical_ndarray = scaler.fit_transform(numerical_not_scaled)
numerical = pd.DataFrame(numerical_ndarray)
numerical
```

Out[21]:

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
|-------|-----------|-----------|-----------|----------|-----------|-----------|----------|-----------|
| 0 | 1.533034 | 0.010471 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| 1 | 1.628993 | -0.421501 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| 2 | -0.290186 | -0.124520 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| 3 | -0.002309 | -0.413787 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| 4 | 1.533034 | 0.187888 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 41183 | 3.164336 | 0.292025 | -0.565922 | 0.195414 | -0.349494 | -0.752343 | 2.058168 | -2.224953 |
| 41184 | 0.573445 | 0.481012 | -0.565922 | 0.195414 | -0.349494 | -0.752343 | 2.058168 | -2.224953 |
| 41185 | 1.533034 | -0.267225 | -0.204909 | 0.195414 | -0.349494 | -0.752343 | 2.058168 | -2.224953 |
| 41186 | 0.381527 | 0.708569 | -0.565922 | 0.195414 | -0.349494 | -0.752343 | 2.058168 | -2.224953 |
| 41187 | 3.260295 | -0.074380 | 0.156105 | 0.195414 | 1.671136 | -0.752343 | 2.058168 | -2.224953 |

41188 rows × 10 columns



In [22]:

```
# Riunisco i dati numerici riscaldati con la parte codificata delle variabili categoriali
features=numerical.join(categorical_encoded, lsuffix='_caller', rsuffix='_other')
features
```

Out[22]:

| | 0_caller | 1_caller | 2_caller | 3_caller | 4_caller | 5_caller | 6_caller | 7_caller |
|-------|-----------|-----------|-----------|----------|-----------|-----------|----------|-----------|
| 0 | 1.533034 | 0.010471 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| 1 | 1.628993 | -0.421501 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| 2 | -0.290186 | -0.124520 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| 3 | -0.002309 | -0.413787 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| 4 | 1.533034 | 0.187888 | -0.565922 | 0.195414 | -0.349494 | 0.648092 | 0.722722 | 0.886447 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 41183 | 3.164336 | 0.292025 | -0.565922 | 0.195414 | -0.349494 | -0.752343 | 2.058168 | -2.224953 |
| 41184 | 0.573445 | 0.481012 | -0.565922 | 0.195414 | -0.349494 | -0.752343 | 2.058168 | -2.224953 |
| 41185 | 1.533034 | -0.267225 | -0.204909 | 0.195414 | -0.349494 | -0.752343 | 2.058168 | -2.224953 |
| 41186 | 0.381527 | 0.708569 | -0.565922 | 0.195414 | -0.349494 | -0.752343 | 2.058168 | -2.224953 |
| 41187 | 3.260295 | -0.074380 | 0.156105 | 0.195414 | 1.671136 | -0.752343 | 2.058168 | -2.224953 |

41188 rows × 65 columns

In [23]:

```
# Trasformo il DataFrame in un ndarray
features_array = features.to_numpy()
features_array
```

Out[23]:


```
array([[ 1.53303429,  0.01047142, -0.56592197, ...,  0.          ,
         1.          ,  0.          ],
       [ 1.62899323, -0.42150051, -0.56592197, ...,  0.          ,
         1.          ,  0.          ],
       [-0.29018564, -0.12451981, -0.56592197, ...,  0.          ,
         1.          ,  0.          ],
       ...,
       [ 1.53303429, -0.26722482, -0.20490853, ...,  0.          ,
         1.          ,  0.          ],
       [ 0.38152696,  0.70856893, -0.56592197, ...,  0.          ,
         0.          ,  1.          ],
       [ 3.26029527, -0.07438021,  0.15610492, ...,  0.          ,
         1.          ,  0.          ]])
```

In [24]:

```
# Riporto di nuovo il grafico a torta per la label questa volta codificate con LinearEncoder
labels_le = labels.apply(le.fit_transform)
labels_le.y.value_counts().plot(kind='pie', subplots=True, startangle=90,
figsize=(10,5), autopct='%1.1f%%')
```

Out[24]:

```
array([<matplotlib.axes._subplots.AxesSubplot object at 0x7f94b810ed
f0>],
      dtype=object)
```



In [25]:

```
# Divido il dataset nelle istanze di training e di test
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(features_array, labels_le, test_size=0.5)
```

ADASYN Oversampling

In [26]:

```
# SOLO sulle istanze di training agirò con un algoritmo di Oversampling, ADASYN
from imblearn.over_sampling import ADASYN
oversample = ADASYN()
```

In [27]:


```
# Resample
X_train, y_train = oversample.fit_resample(X_train, y_train)
```

In [28]:

```
# Come è cambiata la distribuzione?  
y_train.y.value_counts().plot(kind='pie', subplots=True, startangle=90,  
figsize=(10,5), autopct='%1.1f%%')
```

Out[28]:

```
array([<matplotlib.axes._subplots.AxesSubplot object at 0x7f94b09a54  
60>],  
      dtype=object)
```




In [29]:

```
# Sul test set le proporzioni sono quelle originarie...  
y_test.y.value_counts().plot(kind='pie', subplots=True, startangle=90,  
figsize=(10,5), autopct='%1.1f%%')
```

Out[29]:

```
array([<matplotlib.axes._subplots.AxesSubplot object at 0x7f94b096d5  
50>],  
      dtype=object)
```



In [30]:

```
# Porto training set e test set in formato ndarray
y_train = y_train.to_numpy()
y_test = y_test.to_numpy()
print(type(y_train), type(y_test), type(X_train), type(X_test))
```

```
<class 'numpy.ndarray'> <class 'numpy.ndarray'> <class 'numpy.ndarra
y'> <class 'numpy.ndarray'>
```

A Machine is Learning ...

In [31]:

```
# È necessario avere TF >= 2 per utilizzare il pacchetto tensorflow-privacy
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
import tensorflow as tf
print(tf.__version__)
```

2.3.1

In [32]:

```
# Definisco l' architettura di una rete neurale "profonda" che alleneremo
model = Sequential([
    Dense(20, activation='relu', input_shape=(65,)),
    Dropout(.3),
    Dense(20, activation='relu'),
    Dropout(.3),
    Dense(20, activation='relu'),
    Dropout(.3),
    Dense(20, activation='relu'),
    Dense(1, activation='sigmoid'),
])
```

In [33]:

```
# Adam Optimizer con privacy learning
from tensorflow_privacy.privacy.optimizers.dp_optimizer_keras import DPKerasAdam
Optimizer
optimizer = DPKerasAdamOptimizer(
    l2_norm_clip=1.5,
    noise_multiplier=1.3,
    num_microbatches=1033,
    learning_rate=0.025)
```

In [34]:

```
model.compile(
    optimizer=optimizer,
    loss='binary_crossentropy',
    metrics=['accuracy', 'Precision', 'Recall']
)
```

In [35]:

```
history_dpadam=model.fit(
    X_train,
    y_train,
    epochs=10,
    validation_split=0.33,
    batch_size=24,
)
```

Epoch 1/10

```
1016/1016 [=====] - 2s 2ms/step - loss: 0.0
328 - accuracy: 0.9887 - precision: 0.9840 - recall: 0.9709 - val_lo
ss: 8.6755e-07 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_
recall: 1.0000
```

Epoch 2/10

```
1016/1016 [=====] - 2s 2ms/step - loss: 0.0
237 - accuracy: 0.9954 - precision: 0.9888 - recall: 0.9932 - val_lo
ss: 7.6175e-10 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_
recall: 1.0000
```

Epoch 3/10

```
1016/1016 [=====] - 2s 1ms/step - loss: 0.0
025 - accuracy: 0.9994 - precision: 0.9982 - recall: 0.9993 - val_lo
ss: 1.9401e-12 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_
recall: 1.0000
```

Epoch 4/10

```
1016/1016 [=====] - 2s 2ms/step - loss: 0.0
348 - accuracy: 0.9975 - precision: 0.9926 - recall: 0.9974 - val_lo
ss: 0.0016 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_reca
ll: 1.0000
```

Epoch 5/10

```
1016/1016 [=====] - 1s 1ms/step - loss: 0.0
652 - accuracy: 0.9867 - precision: 0.9531 - recall: 0.9966 - val_lo
ss: 5.1184e-06 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_
recall: 1.0000
```

Epoch 6/10

```
1016/1016 [=====] - 1s 1ms/step - loss: 0.0
130 - accuracy: 0.9960 - precision: 0.9849 - recall: 0.9993 - val_lo
ss: 7.6155e-15 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_
recall: 1.0000
```

Epoch 7/10

```
1016/1016 [=====] - 1s 1ms/step - loss: 0.0
106 - accuracy: 0.9982 - precision: 0.9937 - recall: 0.9992 - val_lo
ss: 6.2666e-19 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_
recall: 1.0000
```

Epoch 8/10

```
1016/1016 [=====] - 1s 1ms/step - loss: 0.0
090 - accuracy: 0.9973 - precision: 0.9900 - recall: 0.9992 - val_lo
ss: 5.7273e-10 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_
recall: 1.0000
```

Epoch 9/10

```
1016/1016 [=====] - 1s 1ms/step - loss: 0.0
046 - accuracy: 0.9988 - precision: 0.9955 - recall: 0.9997 - val_lo
ss: 1.0318e-23 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_
recall: 1.0000
```

Epoch 10/10

```
1016/1016 [=====] - 2s 1ms/step - loss: 0.0
075 - accuracy: 0.9984 - precision: 0.9940 - recall: 0.9995 - val_lo
ss: 7.0661e-24 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_
recall: 1.0000
```

In [36]:

```
print("Valuto il modello sui dati di test")
results = model.evaluate(X_test, y_test, batch_size=48)
print("test loss, test acc:", results)

print("Genero 10 predizioni")
predictions = model.predict(X_test[:10])
print("predizioni: \n", predictions)
y_test[:10]
```

Valuto il modello sui dati di test

430/430 [=====] - 0s 804us/step - loss: 2.9

926e-20 - accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000

test loss, test acc: [2.9925607722798456e-20, 1.0, 1.0, 1.0]

Genero 10 predizioni

predizioni:

```
[[0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [1.]
 [0.]]
```

Out[36]:

```
array([[0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [1],
       [0]])
```

Sui dati di test il modello raggiunge accuracy, precision, recall 1. È incoraggiante ma apre domande sulla correttezza metodologica della ricerca.

In [37]:

```
# Adam Optimizer con privacy learning
model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy', 'Precision', 'Recall']
)
```

In [38]:

```
history_adam = model.fit(  
    X_train,  
    y_train,  
    epochs=10,  
    validation_split=0.33,  
    batch_size=24,  
)
```

```
Epoch 1/10
1016/1016 [=====] - 2s 2ms/step - loss: 0.0065 - accuracy: 0.9977 - precision: 0.9910 - recall: 0.9998 - val_loss: 4.6750e-32 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
Epoch 2/10
1016/1016 [=====] - 1s 1ms/step - loss: 0.0094 - accuracy: 0.9982 - precision: 0.9935 - recall: 0.9995 - val_loss: 3.9171e-31 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
Epoch 3/10
1016/1016 [=====] - 1s 1ms/step - loss: 0.0047 - accuracy: 0.9984 - precision: 0.9935 - recall: 1.0000 - val_loss: 1.5085e-32 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
Epoch 4/10
1016/1016 [=====] - 1s 1ms/step - loss: 0.0041 - accuracy: 0.9989 - precision: 0.9956 - recall: 0.9998 - val_loss: 1.0664e-36 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
Epoch 5/10
1016/1016 [=====] - 2s 1ms/step - loss: 0.0043 - accuracy: 0.9986 - precision: 0.9947 - recall: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
Epoch 6/10
1016/1016 [=====] - 2s 2ms/step - loss: 0.0053 - accuracy: 0.9983 - precision: 0.9934 - recall: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
Epoch 7/10
1016/1016 [=====] - 2s 2ms/step - loss: 0.0026 - accuracy: 0.9993 - precision: 0.9972 - recall: 1.0000 - val_loss: 0.0000e+00 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
Epoch 8/10
1016/1016 [=====] - 2s 2ms/step - loss: 0.0082 - accuracy: 0.9990 - precision: 0.9963 - recall: 0.9998 - val_loss: 0.0000e+00 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
Epoch 9/10
1016/1016 [=====] - 2s 2ms/step - loss: 0.0033 - accuracy: 0.9989 - precision: 0.9960 - recall: 0.9998 - val_loss: 0.0000e+00 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
Epoch 10/10
1016/1016 [=====] - 2s 2ms/step - loss: 0.0115 - accuracy: 0.9992 - precision: 0.9971 - recall: 0.9997 - val_loss: 0.0000e+00 - val_accuracy: 1.0000 - val_precision: 1.0000 - val_recall: 1.0000
```

In [39]:

```
print("Valuto il modello sui dati di test")
results = model.evaluate(X_test, y_test, batch_size=48)
print("test loss, test acc:", results)

print("Genero 10 predizioni")
predictions = model.predict(X_test[:10])
print("predizioni: \n", predictions)
y_test[:10]
```

Valuto il modello sui dati di test

430/430 [=====] - 0s 803us/step - loss: 7.9

621e-37 - accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000

test loss, test acc: [7.962126754925632e-37, 1.0, 1.0, 1.0]

Genero 10 predizioni

predizioni:

```
[[0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [0.]
 [1.]
 [0.]]
```

Out[39]:

```
array([[0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [1],
       [0]])
```

Conclusioni

Si sarebbe potuto ipotizzare che usare una versione di ADAM dove SGD è sostituito da DP-SGD potesse impattarne le performance, almeno nel caso in esame ciò non sembra particolarmente vero anche se il valore della funzione di costo dopo all' ultima epoca è più alto.

Riporto il sommario di tutti gli eventi nel training dei due modelli

In [40]:

history_dpadam.history

Out[40]:

```
{'loss': [0.032797202467918396,
0.023743323981761932,
0.0024960124865174294,
0.03479241579771042,
0.06524933874607086,
0.013013748452067375,
0.010589821264147758,
0.009014119394123554,
0.004550587851554155,
0.0075330547988414764],
'accuracy': [0.9886722564697266,
0.9954442977905273,
0.999384343624115,
0.997455358505249,
0.9867432713508606,
0.9959778189659119,
0.99819415807724,
0.9972501397132874,
0.9987687468528748,
0.9983583092689514],
'precision': [0.9840223789215088,
0.9888349771499634,
0.9982143044471741,
0.9925602674484253,
0.9530618786811829,
0.984943151473999,
0.9936964511871338,
0.9900161027908325,
0.9954684972763062,
0.9940197467803955],
'recall': [0.970908522605896,
0.9931740760803223,
0.9993498921394348,
0.997399628162384,
0.9965870380401611,
0.9993498921394348,
0.9991874098777771,
0.9991874098777771,
0.9996749758720398,
0.9995124340057373],
'val_loss': [8.675520462020359e-07,
7.617529385051114e-10,
1.9401481289593736e-12,
0.001569235697388649,
5.118351054989034e-06,
7.61548942261386e-15,
6.266587971803729e-19,
5.727286978007839e-10,
1.031805530454931e-23,
7.06608666330242e-24],
'val_accuracy': [1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0],
'val_precision': [1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0],
'val_recall': [1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0]}
```

In [41]:

```
history_adam.history
```

Out[41]:

```
{'loss': [0.006491903215646744,  
0.009353882633149624,  
0.0046739354729652405,  
0.004108505789190531,  
0.004250656813383102,  
0.0053377775475382805,  
0.0025714579969644547,  
0.008214999921619892,  
0.00332126603461802,  
0.01148801390081644],  
'accuracy': [0.9976605772972107,  
0.9982351660728455,  
0.9983583092689514,  
0.9988508224487305,  
0.9986456036567688,  
0.9983172416687012,  
0.9993022680282593,  
0.9990149736404419,  
0.9989328980445862,  
0.9991791248321533],  
'precision': [0.9909793734550476,  
0.993537962436676,  
0.9935411214828491,  
0.9956303834915161,  
0.9946653842926025,  
0.9933806657791138,  
0.9972447156906128,  
0.9962753057479858,  
0.9959527254104614,  
0.9970821738243103],  
'recall': [0.9998374581336975,  
0.9995124340057373,  
1.0,  
0.9998374581336975,  
1.0,  
1.0,  
1.0,  
1.0,  
0.9998374581336975,  
0.9998374581336975,  
0.9996749758720398],  
'val_loss': [4.675048003206162e-32,  
3.9170506895812443e-31,  
1.5085249777848945e-32,  
1.0664158434296163e-36,  
0.0,  
0.0,  
0.0,  
0.0,  
0.0,  
0.0],  
'val_accuracy': [1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0],  
'val_precision': [1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.  
0],  
'val_recall': [1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0]}
```